WHITE PAPER · UK / EU IP & AI LAW · v1.0 · 2026

Training Data Disclosure in AI Copyright Litigation

The Post-Report Provenance Procedure: a three-stage civil-procedure framework operating through PD 57AD in the Business and Property Courts of England and Wales.

R.A. Aswin Krishna

IP-AI Practitioner · Incoming LLM (Intellectual Property Law), Queen's University Belfast

ABSTRACT

The UK Government's *March 2026 Report on Copyright and Artificial Intelligence* identified an evidentiary asymmetry at the heart of AI-copyright litigation: claimants cannot inspect training corpora, while EU AI Act Article 53 disclosures are aggregated to a level that does not establish causal nexus for individual works. *Getty Images (US) Inc v Stability AI Ltd* [2025] EWHC 2863 (Ch) confirmed the limits of the secondary-infringement route in the United Kingdom. The substantive law has not moved to fill the gap; the procedural toolkit has not been adapted to the asymmetry.

This paper proposes the *Post-Report Provenance Procedure* (PRPP) — a three-stage civil-procedure framework operating through Practice Direction 57AD in the Business and Property Courts. PRPP synthesises three procedural levers that already exist (CPR r.6.37 good-arguable-case; PD 57AD Model C Extended Disclosure; and the Wisniewski adverse-inference doctrine) and applies them to AI-copyright disclosure as a coherent system. PRPP does not change the substantive test for infringement under CDPA 1988 s.16; it changes how the factual element of ingestion can be established in litigation.

Keywords: *AI & copyright; training data disclosure; PD 57AD; CPR Pt 31; Model C Extended Disclosure; Wisniewski; Getty v Stability AI; civil procedure; UK IP law.*

Suggested citation: *R.A. Aswin Krishna, Training Data Disclosure in AI Copyright Litigation: The Post-Report Provenance Procedure (White Paper, 2026). Companion software: Libra Contract Guardian v1.1.*

§ I

Introduction: An Evidentiary Asymmetry

Generative artificial-intelligence systems are trained on corpora the contents of which are, in any meaningful sense, unknowable to the rights-holders whose works they may contain. This asymmetry is not an incidental feature of the technology; it is a defining one. Modern foundation models ingest hundreds of billions of tokens drawn from common-crawl scrapes, licensed corpora, and curated datasets that the developer alone holds in fully-attributed form.

When a copyright owner suspects that her works have been ingested without licence, she encounters a problem that no traditional copyright case has had to confront: she cannot, by any ordinarily-available means, test whether the suspected ingestion in fact occurred. The training corpus is held internally; the model weights are a probabilistic compression of an unknown distribution; the EU AI Act Article 53(1)(d) obligation to publish a 'sufficiently detailed summary' is calibrated to aggregated transparency rather than per-work disclosure.[1]

The UK Government's *March 2026 Report on Copyright and Artificial Intelligence*[2] identified this asymmetry without resolving it. The Report formally abandoned the previously preferred broad TDM exception with opt-out (Option 3), deferred decision on alternative reforms pending further evidence, and expressly preserved the pre-existing balance struck by CDPA 1988 s.29A. What it did not do — and what this paper takes up — is propose a procedural framework by which the asymmetry can be addressed within existing law while the policy debate continues.

The argument advanced here is that the Civil Procedure Rules and Practice Direction 57AD already contain the tools required. They have not been applied to AI-copyright disclosure as a coherent system, but their constituent elements — the good-arguable-case threshold, Model C Extended Disclosure, the confidentiality ring, and the discretionary adverse inference under *Wisniewski v Central Manchester HA* — are doctrinally settled and judicially familiar. Their synthesis into a three-stage procedure is what this paper calls the *Post-Report Provenance Procedure* (PRPP).

§ II

Background: The State of AI-Copyright Litigation

II.1 Getty v Stability AI and the limits of the secondary route

Getty Images (US) Inc v Stability AI Ltd [2025] EWHC 2863 (Ch)[3] is the most consequential decision yet rendered by an English court on the AI-copyright question. The claimant pursued, among other heads, a secondary-infringement claim on the footing that the trained model weights themselves constituted an infringing copy under CDPA 1988 ss.22-23. Joanna Smith J rejected this theory on the evidence: the weights were not, on the technical findings, a recoverable copy of the training images.

The decision is presently under appeal, and a contrary view at appellate level cannot be excluded. But the more important point for present purposes is structural rather than doctrinal: even if the secondary route is reinstated, it does nothing to address the prior, factual question of whether the protected works were in fact ingested. That question remains an evidentiary problem first, and a doctrinal problem second.

II.2 The aggregation problem in EU AI Act Article 53

The EU AI Act, Regulation (EU) 2024/1689, requires providers of general-purpose AI models to publish a 'sufficiently detailed summary' of the content used for training (Art. 53(1)(d)). The implementing template prescribes summary disclosures at the level of dataset categories — public web crawl, licensed text, etc. — rather than at the level of individual works.[4]

From a regulatory-transparency perspective this is intelligible: the Article was drafted to balance copyright awareness against trade-secrecy and competitive concerns. From a litigation-evidence perspective, however, an aggregate disclosure does not establish causal nexus. A claimant cannot prove that *her particular work* was ingested by pointing to a published summary stating that 'public web data' was used — even if her work was undeniably available on the public web. The aggregation level forecloses precisely the inferential step the claimant must make.

II.3 The procedural vacuum

What the rights-holder requires, and what neither the UK March 2026 Report nor the EU AI Act provides, is a procedural mechanism by which she may, on adequate prima facie evidence, compel disclosure of information sufficient to confirm or refute the factual proposition that her works were ingested. PRPP fills that vacuum.

§ III

The PRPP Framework

PRPP is staged. Each stage corresponds to a distinct procedural lever in the existing CPR / PD 57AD architecture, and each must be cleared before the next is reached. The staging is what makes PRPP proportionate: it ensures that costly disclosure is not ordered without antecedent evidence sufficient to justify it.

III.1 Stage One — The Prima Facie Trigger

The claimant must clear a threshold of plausible inference of ingestion, calibrated to the 'good arguable case' standard articulated in CPR r.6.37 and the jurisdictional gateway authorities — not the higher 'real prospect of success' standard under CPR Part 24.[5] Three evidentiary routes are recognised:

(a) Hosted-repository evidence. The claimant adduces factual evidence that her protected portfolio was hosted on a domain comprehensively scraped by a known dataset (Common Crawl, LAION-5B, Books3 are the most common candidates), and cross-references this with the developer's public model-card disclosures. This route is the strongest because it is objective and publicly verifiable.

(b) Circumstantial regurgitation. The model, under targeted prompting, reproduces protected content near-verbatim. The Stanford-Yale study by Ahmed et al. (2026) provides the methodology and a calibration baseline: nv-recall of 95.8% on protected books for Claude 3.7 Sonnet under BoN=258 jailbreak conditions; 76.8% on Harry Potter content for Gemini 2.5 Pro without jailbreaking.[6] This is forensic evidence in the black-box model: it does not require access to weights or corpus.

(c) Membership Inference Attack (MIA). A statistical methodology, deployed via CPR Part 35 expert evidence, that queries model outputs to determine — with probabilistic confidence — whether a specific data point was in the training set. White-box MIA (full model access) is highly reliable but unavailable to the claimant; black-box variants are admissible as a statistical indicator but not as definitive proof. PRPP treats Stage One MIA as supportive evidence, not as a stand-alone trigger.

III.2 Stage Two — Model C Extended Disclosure

Once Stage One is cleared, the claimant seeks Model C Extended Disclosure under PD 57AD of narrow classes of material: cryptographic hash manifests (SHA-256) and deduplication logs of the training corpus.[7] The disclosure is narrow because the question is narrow: did the developer's training set contain a hash matching the claimant's work? A binary match — established in seconds by a jointly-instructed expert — answers it.

Production occurs within a confidentiality ring. The choice of tier is fact-sensitive:

- **External-eyes-only**, per *IPCom v HTC Europe* [2013] EWHC 2880 (Ch),[8] where trade-secrecy concerns predominate and the developer is well-resourced.
- **Three-tier**, per *Mitsubishi Electric v OnePlus* [2020] EWCA Civ 1562, for proprietary-algorithm cases where internal counsel involvement is appropriate.
- **Self-certified (SME)**, where the developer is a small or medium enterprise. The March 2026 Report at p.66 expressly identified disproportionate cost as a barrier to innovation; PRPP responds with a lower-cost tier that preserves the procedural architecture without imposing external-counsel costs.

Proportionality is governed by PD 57AD para 6.4 and CPR r.1.1. The narrowness of the disclosure — hash manifests, not full corpora — is what makes the procedure proportionate even in lower-value claims.

III.3 Stage Three — Adverse Inference

Where the developer fails to preserve or to produce hash records without credible non-culpable explanation, the court may exercise its discretion to draw an adverse inference on the factual issue of ingestion. The foundational principle is set out in *Wisniewski v Central Manchester Health Authority* [1998] EWCA Civ 596; [1998] PIQR P324 (Brooke LJ),[9] with the doctrinal extension to absent documents articulated in *Wetton v Ahmed* [2011] EWCA Civ 610 (Arden LJ at [14]) and applied to electronic records in *Earles v Barclays Bank Plc* [2009] EWHC 2500 (Mercantile).

PRPP requires the court at Stage Three to distinguish two situations, consistent with the Earles principle that the duty to preserve engages once proceedings are contemplated, not before:

• **Deliberate spoliation**: non-retention *after* formal notice of claim, or refusal to comply with a disclosure order without credible justification. The Wisniewski / Wetton inference is appropriate.
• **Routine data-minimisation**: pre-litigation overwrites in the ordinary course of business, server-optimisation deletions, and data-protection-compliant minimisation regimes adopted before any knowledge of contemplated proceedings. The inference should not be drawn from these alone — Earles confirms there is no general pre-action preservation duty.

The discretion is not mandatory and supports the factual finding of ingestion; it does not discharge the claimant's ultimate burden on the substantive infringement question. PRPP is, throughout, a procedural mechanism — not a substantive doctrine.

§ IV

Anticipated Objections

IV.1 'This is disclosure overreach.'

PRPP does not order disclosure of the training corpus. It orders disclosure of cryptographic hashes — fixed-length strings that reveal nothing about the underlying content other than whether a specific known input was present. The mechanism is more analogous to a forensic hash-comparison in computer-misuse litigation than to a corpus-wide data-room exercise.

IV.2 'It will chill UK AI investment.'

The opposite. The status quo — under which UK courts have no settled procedure for handling training-data evidence — produces uncertainty for both rights-holders and developers. PRPP creates a defined procedural pathway with proportionality safeguards, which is precisely what well-advised developers prefer to ad-hoc disclosure battles. The SME tier is designed to ensure smaller developers are not priced out.

IV.3 'The substantive law has not been settled.'

Correct, and PRPP does not require it to be. The framework operates at the procedural layer: it allows the factual question of ingestion to be answered. Whether that answer entails infringement is a separate question, governed by the existing substantive tests in CDPA 1988 ss.16-17 and the leading cases (*Designers Guild v Russell Williams*; *Infopaq*). The two questions are logically separable; PRPP separates them.

§ V

Practical Implications

PRPP has implications for three constituencies.

For rights-holders and their advisers, PRPP provides a pleadable procedural pathway. A claim particularised under PRPP should: (i) plead Stage One evidence in the Particulars of Claim; (ii) include 'algorithmic ingestion' as a contested Issue for Disclosure on the Disclosure Review Document; (iii) seek a Stage Two Model C order at the first Case Management Conference; and (iv) preserve the Stage Three argument by sending a pre-action preservation letter at the earliest opportunity.

For AI developers and in-house counsel, PRPP creates an incentive to maintain hash-manifest records of training corpora as a matter of good practice. Such records are inexpensive to create, do not compromise trade secrecy (a hash reveals nothing about content), and substantially reduce the developer's exposure to a Stage Three adverse inference if they are produced when called for.

For courts and the Disclosure Working Group, PRPP offers a worked-through application of PD 57AD to a problem the Practice Direction was not specifically designed for but accommodates naturally. The framework requires no rule-change; it requires only the recognition that 'algorithmic ingestion' is a discrete Issue for Disclosure susceptible to Model C treatment, and that hash-manifest records are a discrete class of documents within the meaning of PD 57AD para 6.5.

§ VI

Conclusion

The argument of this paper is conservative in form and ambitious only in its synthesis. The CPR r.6.37 threshold has been settled for two decades. PD 57AD has been operating in the Business and Property Courts since 2019. The *Wisniewski* doctrine is older still. What is new is the proposition that these three settled procedural tools, applied in sequence to the discrete problem of AI-training-data disclosure, constitute a coherent and proportionate response to the evidentiary asymmetry the March 2026 Report identified.

The Post-Report Provenance Procedure does not change what counts as infringement. It changes how the factual element of ingestion can be established. That is sufficient. The substantive questions — what constitutes a substantial part; whether transient copies engage s.17; the proper construction of CDPA 1988 s.29A in commercial settings — remain open, as they should. PRPP exists to ensure those questions can be reached, on a properly-developed factual record, in the courts of England and Wales.

§ Endnotes

1. Regulation (EU) 2024/1689 (the EU AI Act), Article 53(1)(d). The European AI Office's implementing template (2025) calibrates disclosure to dataset-category level.

2. UK Government, *Report on Copyright and Artificial Intelligence* (March 2026), available at gov.uk.

3. *Getty Images (US) Inc v Stability AI Ltd* [2025] EWHC 2863 (Ch). The judgment of Smith J holds that, on the technical evidence adduced, the trained model weights did not store recoverable copies of the training images; the claimant's secondary-infringement claim under CDPA 1988 ss.22-23 accordingly failed. Appeal pending.

4. European Commission, *Template for the Public Summary of Training Content for General-Purpose AI Models*, adopted by the AI Office on 24 July 2025, in force for new models from 2 August 2025. The template is calibrated to dataset-category, source-domain, and approximate-volume disclosure, rather than per-work attribution. See also EU AI Act Code of Practice on Transparency, Copyright, and Safety (10 July 2025).

5. The good-arguable-case standard is articulated in *Brownlie v Four Seasons Holdings Inc* [2017] UKSC 80 and is the relevant threshold under CPR r.6.37 for jurisdictional gateways. It is lower than the CPR Part 24 standard and is the appropriate calibration for a Stage One PRPP trigger. For Disclosure under PD 57AD the threshold is gating-by-Issue rather than gating-by-merits, but the structural analogy holds.

6. Ahmed Ahmed et al., *Extracting books from production language models* (2026) arXiv:2601.02671 [cs.CL]. **Preprint, not peer-reviewed at time of writing**; institutional affiliations as stated by the authors. Headline figures: Claude 3.7 Sonnet at 95.8% nv-recall under Best-of-N jailbreak (BoN=258); Gemini 2.5 Pro at 76.8% on Harry Potter and the Sorcerer's Stone without jailbreaking. The methodology is black-box and reproducible via prompt-only access. Cited here as preliminary empirical evidence of verbatim-memorisation capacity in production models, not as proof of infringement.

7. PD 57AD Model C is the 'request-led search-based' disclosure model, in which parties identify narrow classes of documents tied to the contested Issues for Disclosure. It is well-suited to AI-training-data disclosure precisely because the relevant class (hash manifests of ingested content) is narrow and defined.

8. *IPCom GmbH & Co KG v HTC Europe Co Ltd* [2013] EWHC 2880 (Ch) at [360]–[377] (Roth J), establishing a tightly-drawn confidentiality ring restricting source-code inspection to external lawyers and independent experts. *Mitsubishi Electric Corporation v*

OnePlus Technology (Shenzhen) Co Ltd [2020] EWCA Civ 1562 at [18]–[23] upholds three-tier rings for proprietary algorithms; see also *TQ Delta v Zyxel* [2018] EWHC 1515 (Ch) (external-eyes-only is exceptional) and *Infederation Ltd v Google LLC* [2020] EWHC 657 (Ch) at [27]–[42] (Roth J) on the principles applicable to confidentiality rings.

9. *Wisniewski v Central Manchester Health Authority* [1998] EWCA Civ 596; [1998] PIQR P324 (Brooke LJ at p340) — the foundational principles for adverse inferences from absent witnesses. The doctrinal extension to absent documents was articulated in *Wetton v Ahmed* [2011] EWCA Civ 610 (Arden LJ at [14]); see also *Earles v Barclays Bank Plc* [2009] EWHC 2500 (Mercantile) (electronic records; preservation duty engages once proceedings are contemplated). The inference is discretionary, not mandatory: see *Magdeev v Tsvetkov* [2020] EWHC 887 (Comm) (Cockerill J) for the modern restatement of the structured test.

Companion software: *Libra Contract Guardian* v1.1 (2026) operationalises the PRPP framework as a Streamlit demonstrator, with three engines (PRPP, TDM, Playbook) running against a curated UK/EU IP-AI knowledge base. The Playbook engine implements TF-IDF vector-space cosine similarity against gold-standard contract corpora; verification harness in *test_playbook.py*. Source available on request.

Author contact: R.A. Aswin Krishna · aswinkrishna2277@gmail.com · Incoming LLM (Intellectual Property Law), Queen's University Belfast.